
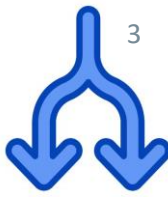


Practical Concurrent and Parallel Programming VI Performance and Scalability

Jørgen Staunstrup



- Performance versus scalability
- Scalability, speed-up and loss classification
Example: QuickSort
- Executors and Future
Example: count Prime Factors
- Hash maps, a scalability case study



Week 3

Speedup for quicksort:

3.6 using 8 threads

2.9 using 4 threads

Speedup for counting primes:

3.9 using 8 threads

2.3 using 4 threads



- **Performance versus scalability**
- Scalability, speed-up and loss classification
Example: QuickSort
- Executors and Future
Example: count Prime Factors
- Hash maps, a scalability case study

Performance and scalability



Performance (of software)

- **Latency:** time till first result (response time)
- **Throughput:** results per second

Scalability (one way to improve performance)


Improve throughput/latency by adding more resources

One may sacrifice performance for scalability

Maybe OK to be slower on 1 core, if faster on 2 or 4 or ...

Goetz chapter 11

What limits performance?



Suggestions?

What limits performance?



CPU-bound

- Eg. counting prime numbers
- To speed up, add more CPUs (cores) (*exploitation*)

Input/output-bound

- Eg. reading from network
- To speed up, use more tasks (*inherent*)

Synchronization-bound

- Eg. Algorithm using shared data structure
- **To speed up, improve shared data structure**
(*Much of this lecture*)



- Performance versus scalability
- **Scalability, speed-up and loss classification**
Example: QuickSort
- Executors and Future
Example: count primes
- Hash maps, a scalability case study

Quicksort



9

1 2 43 78 19 54 33 21 64 52 17 53

1 2 43 78 19 54 33 21 64 52 17 53

1 2 43 78 19 54 33 21 64 52 17 53



1 2 17 78 19 54 33 21 64 52 43 53



1 2 17 78 19 54 33 21 64 52 43 53




...

1 2 17 21 19 33 54 78 64 52 43 53

1 2 17 21 19 33 54 78 64 52 43 53

Two parts can be
sorted independently

Multi-threaded version of Quicksort



```
class Problem {  
    public int[] arr;  
    public int low, high;  
    ...  
}
```

```
class ProblemHeap {  
    list<Problem> heap= new List<Problem>;  
}
```

```
private static void qsort(Problem problem, ProblemHeap heap) {  
    int[] arr= problem.arr;  
    int a= problem.low;  
    int b= problem.high;  
    ...  
    heap.add(new Problem(arr, a, j)); //qsort(arr, a, j);  
    heap.add(new Problem(arr, i, b)); //qsort(arr, i, b);  
}
```

Java code for Quicksort (1)



```
public static void problemHeapStart(int threadCount, int pSize, int[] intArray) {
    ProblemHeap heap= new ProblemHeap(1);
    heap.add(new Problem(intArray, 0, pSize-1));

    for (int t=0; t<threadCount; t++) {
        threads[t]= new Thread( () -> { try {
            Problem newProblem= heap.getProblem();
            while (newProblem != null) { // when newProblem == null alg stops
                qsort(newProblem, heap);
                newProblem= heap.getProblem();
            }
        } catch (InterruptedException exn) { } //needed because getProblem may wait
        });
    }
}
```

Java code for Quicksort (2)

12



We use Mark8Setup to measure runtime


```
Benchmark.Mark8Setup("Problem heap quicksort",
    String.format("%2d", threadCount),
    new Benchmarkable() {
        public void setup() {
            SearchAndSort.shuffle(intArray);
            problemHeapStart(threadCount, pSize, intArray);
        }
        public double applyAsDouble(int i) {
            problemHeapFinish(threadCount, intArray); return 0.0;
        }
    }
);
```

Code in: **Week06: exercises/**
ProblemHeapSortingBenchmarkable.java



Motivation

- Threads are expensive to start - executors reuse threads
- Problem heap - breaking a problem down to smaller problems (tasks)



Task producers, and *Workers* are threads

Workers may themselves produce new tasks

The Task pool and Worker pool together is called an *Executor service*



```
new Thread(runnable1).start();  
...  
new Thread(runnable2).start();  
...  
new Thread(runnable3).start();
```

Threads are expensive !

```
ExecutorService pool;
```

```
pool.execute(runnable1);  
...  
pool.execute(runnable2);  
...  
pool.execute(runnable2);
```

Reuse of threads

<https://howtodoinjava.com/java/multi-threading/java-fixed-size-thread-pool-executor-example/>

Quicksort example

16



```
class runProblemHeap extends Thread {
    private final ExecutorService pool;
    public runProblemHeap(...) { pool = Executors.newFixedThreadPool(poolSize); }

    @Override
    public void run() {
        pool.execute( new solveProblem( new Problem(...), pool, c) );
    }
}

class solveProblem implements Runnable {
    ..
    @Override
    public void run() { qsort(... ); }

    // Quicksort
    public static void qsort(..., ExecutorService pool) { ... }
}
```

Quicksort using Executors

17



```
class solveProblem implements Runnable {
    private ExecutorService pool;
    private static int threshold;
    public solveProblem(ExecutorService pool, int threshold, ...) { }
    @Override
    public void run() { qsort(pool, ... ); }
    public static void qsort(ExecutorService pool, ... ) {
        if (a < b) { ... }

        if ((j-a)>=threshold) pool.execute(new solveProblem(pool,threshold));
        else { SearchAndSort.qsort(..., a, j); }

        if ((b-i)>= threshold) pool.execute(new solveProblem(pool,threshold));
        else { SearchAndSort.qsort(..., i, b); }
    }
}
}
```

Code in Week06: Exercises [.../QuicksortExecutor.java](#)

Performance of Executor Quicksort



Executor quicksort	1	98003405.0 ns
Executor quicksort	2	53568593.9 ns
Executor quicksort	4	36397241.3 ns
Executor quicksort	8	21714103.7 ns
Executor quicksort	16	22237307.4 ns
Executor quicksort	32	22510681.9 ns

Speedup = 4.5

A bit better than using native Threads (slide 3)

Quicksort

19



1 2 43 78 19 54 33 21 64 52 17 53

1 2 43 78 19 54 33 21 64 52 17 53

1 2 43 78 19 54 33 21 64 52 17 53



1 2 17 78 19 54 33 21 64 52 43 53



1 2 17 78 19 54 33 21 64 52 43 53



...

1 2 17 21 19 54 33 78 64 52 43 53

1 2 17 21 19 33 54 78 64 52 43 53

Two parts can be
sorted independently

What limits scalability?



Example: growing a crop

- 4 months growth + 1 month harvest if done by 1 person
- Growth (sequential) cannot be speeded up
- Using 30 people to harvest, takes $1/30$ month = 1 day
- Speed-up using many harvesters: $5/(4+1/30) = 1.24$ times faster

Amdahl's law (Goetz 11.2)

F = sequential fraction of problem = $4/5 = 0.8$

N = number of threads (people) = 30

Speed-up $\leq 1/(F+(1-F)/N) = 1/(0.8+0.2/30) = 1.24$

Other types of "loss" limiting scalability



Starvation loss

Separation loss (best threshold)

Saturation loss (locking common data structure)

Braking loss

Møller-Nielsen, P and Staunstrup, J, Problem-heap. A paradigm for multiprocessor algorithms. *Parallel Computing*, 4:63-74, 1987



- Performance versus scalability
- Scalability, speed-up and loss classification
Example: QuickSort
- **Executors and Future**
Example: count Prime Factors
- Hash maps, a scalability case study

Termination in Executors

23



```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

...

```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

```
//main (thread)  
  setUpQS( ... ) // creates thread  
  finishQS( ... ) //wait for all threads to finish  
}
```

Need for a signal from the threads to the main thread !!!

The ExecutorService has methods to help in shutting down.

```
// Executor body  
...  
...
```

```
pool.shutdown();
```

The challenge is: when to shut down?

In the Quicksort example a counter (=size of problem heap) was used to determine when to shutdown

Solution 1: shared counter



```
class countProblems{
    private int c= 1; // counting active threads + problems in heap
    synchronized void incr() { c++; }
    synchronized void decr() { c--; }
    synchronized void reset() { c= 1; }
    synchronized boolean isZero() { return c==0; }

    // The semaphore finished signals termination to main thread
    public Semaphore finished= new Semaphore(0);
}
```




```
public static void qsort(..., countProblems c) {  
    if (a < b) {  
        ...  
        if ((j-a)>= threshold) {  
            c.incr();  
            pool.execute(new solveProblem(new Problem(arr, a, j), pool, c) );  
        };  
        if ((b-i)>= 1000) {  
            c.incr();  
            pool.execute(new solveProblem(new Problem(arr, i, b), pool, c) );  
        }  
        c.decr(); // problem solved, decrement c  
  
        if ( c.isZero() ) { signal termination pool.shutdown(); }  
    }  
}
```

How can we be sure that no other threads are active?

Termination in Executors

27



```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

...

```
class solveProblem implements Runnable {  
    public void run() { qsort(p, ...); }  
}
```

```
//main (thread)  
    setUpQS( ... ) // creates thread  
    finishQS( ... ) //wait for all threads to finish  
}
```

```
private static void finishQS(countProblems c) {  
    try { c.finished.acquire(); }  
    catch (java.lang.InterruptedException e) { ... }; }  
}
```

Solution 2: Future



```
T2:
public Future<Integer> calculate(Integer input) {
    return executor.submit(() -> {
        ... // compute result
        return result;
    });
}

T1:
Future<Integer> future = T2.calculate( ); // start
...
future.get();
```



```
private ExecutorService executor
    = Executors.newSingleThreadExecutor();

public Future<Integer> calculate(Integer input) {
    return executor.submit( //Callable
        () -> {... return ... ; }
    );
}
```

Code in Week06: **LectureCode** [.../futureExample.java](#)

Both are used to specify the code of a thread.

- Runnable cannot return a result
- Callable returns a result (via a Future)

As illustrated by the Quicksort example, Runnables may use shared data (e.g., to deliver a result)

Futures are an example of message passing



```
private static long countParallelN(int range, int taskCount) {
    List<Callable<Long>> tasks= new ArrayList<Callable<Long>>();

    for (int t= 0; t<taskCount; t++) {
        final int from= ...,
                to= ...;
        tasks.add(() -> {
            long count = 0; // Task-local counter
            for (int i=from; i<to; i++) if (isPrime(i)) count++;
            return count;
        });
    }

    long result = 0;
    try {
        List<Future<Long>> futures = executor.invokeAll(tasks);
        for (Future<Long> fut : futures)
            result += fut.get();
    } catch ...
    return result;
}
```

- Performance versus scalability
- Scalability, speed-up and loss classification
Example: QuickSort
- Executors and Future
Example: count Prime Factors
- **Hash maps, a scalability case study**

Scalability of Java Collections



A *collection* is simply an object that groups multiple elements into a single unit

Package: `java.util`

Examples: `ArrayList`, `HashMap`, `TreeSet`, ...

<https://docs.oracle.com/javase/tutorial/collections/intro/index.html>

Methods: `add`, `remove`, `size`, `contains`, ...

Many of the classes have synchronized/concurrent implementations

<https://www.baeldung.com/java-synchronized-collections>

Example: synchronized ArrayList

34



```
import java.util.*;

public class syncCollectionExample {
    public static void main(String[] args) { new syncCollectionExample(); }

    public String getLast(ArrayList<String> l) {
        int last= l.size()-1;
        return l.get(last);
    }

    public static void delete(ArrayList<String> l) {
        int last= l.size()-1;
        l.remove(last);
    }

    public syncCollectionExample() {
        ArrayList<String> a= new ArrayList<String>();
        a.add("A"); ...

        Collection<String> synColl = Collections.synchronizedCollection(a);
        ...
    }
}
```

Goetz p. 80

It is very important to note that:

*Program where all classes are thread-safe \nRightarrow
thread-safe program*

Making the synchronized ArrayList thread safe



36

```
import java.util.*;
public class syncCollectionExample {
    public static void main(String[] args) { new syncCollectionExample(); }

    public String getLast(ArrayList<String> l) {
        synchronized(l) {
            int last= l.size()-1;
            return l.get(last);
        }
    }

    public static void delete(ArrayList<String> l) {
        synchronized(l) {
            int last= l.size()-1;
            l.remove(last);
        }
    }

    public syncCollectionExample() {
        ...
    }
}
```

Goetz p. 80

What if the data structure is huge?



and used by many threads?

for example:

a bank

Facebook updates

...

Would not work if everything I "synchronized"

What can we do?

Reduce locking !!

Example: A huge HashMap

Key value pairs: $\langle k_1, v_1 \rangle, \langle k_2, v_2 \rangle, \dots$


```
class HashMap<K,V> {  
    ... // datastructure  
    public V get(K k) { ... }  
    public V put(K k, V v) { ... }  
    public boolean containsKey(K k) { ... }  
    public int size() { return cachedSize; }  
    public V remove(K k) { ... }  
    ...  
}
```

How to make it thread-safe?

Key	Value
Peter	20487612
Anna	51251218
Lena	34458318
Holger	89545010
Lisa	94959500



Scaling a HashMap




HashMap implementation


35




buckets




ItemNode



ItemNode



ItemNode



Example **get(10406)**

key k is 10406

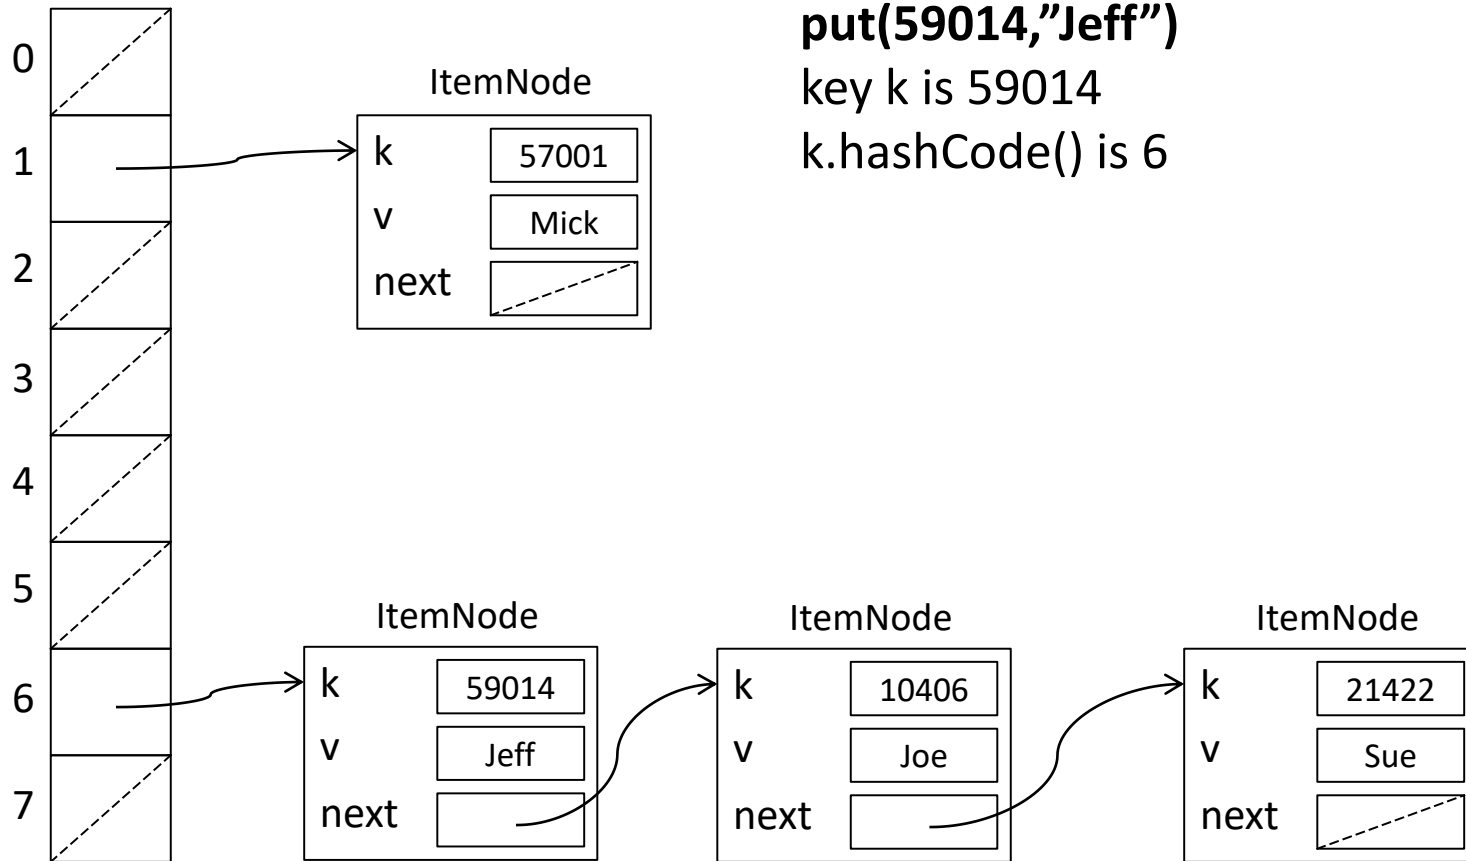
k.hashCode() is 6

HaspMap put

36



buckets



Synchronized implementation




```
static class ItemNode<K,V> {  
    private final K k;  
    private V v;  
    private ItemNode<K,V> next;  
    public ItemNode(K k, V v, ItemNode<K,V> next) { ... }  
}
```

```
class SynchronizedMap<K,V> {  
    private ItemNode<K,V>[] buckets; // guarded by this  
    private int cachedSize;          // guarded by this  
    public synchronized V get(K k) { ... }  
    public synchronized boolean containsKey(K k) { ... }  
    public synchronized int size() { return cachedSize; }  
    public synchronized V put(K k, V v) { ... }  
    public synchronized V remove(K k) { ... }  
}
```

Improving scalability




- Guarding the table with a single lock works
 - ... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
 - use fewer, to illustrate we use 4, locks
 - guard every 4th bucket with the same lock
 - locks[0] guards bucket 0, 4, 8, ... (stripe 0)
 - locks[1] guards bucket 1, 5, 9, ... (stripe 1) et
 - With high probability
 - two operations will work on different stripes
 - hence will take different locks
- Less lock contention, better scalability




Bucket idea



buckets




ItemNode




Locking one will not lock the other

In different stripes


ItemNode




ItemNode



ItemNode



Reducing locking





A web-shop, Facebook, ...

We must give up thread safety,

but still maintain some sort of consistency

Week 13